



# SBI HACKATHON 2026: CONTRIBUTOR TRACK

---

**LOCATION**

Grenoble, France

**DATE**

January 21-23, 2026

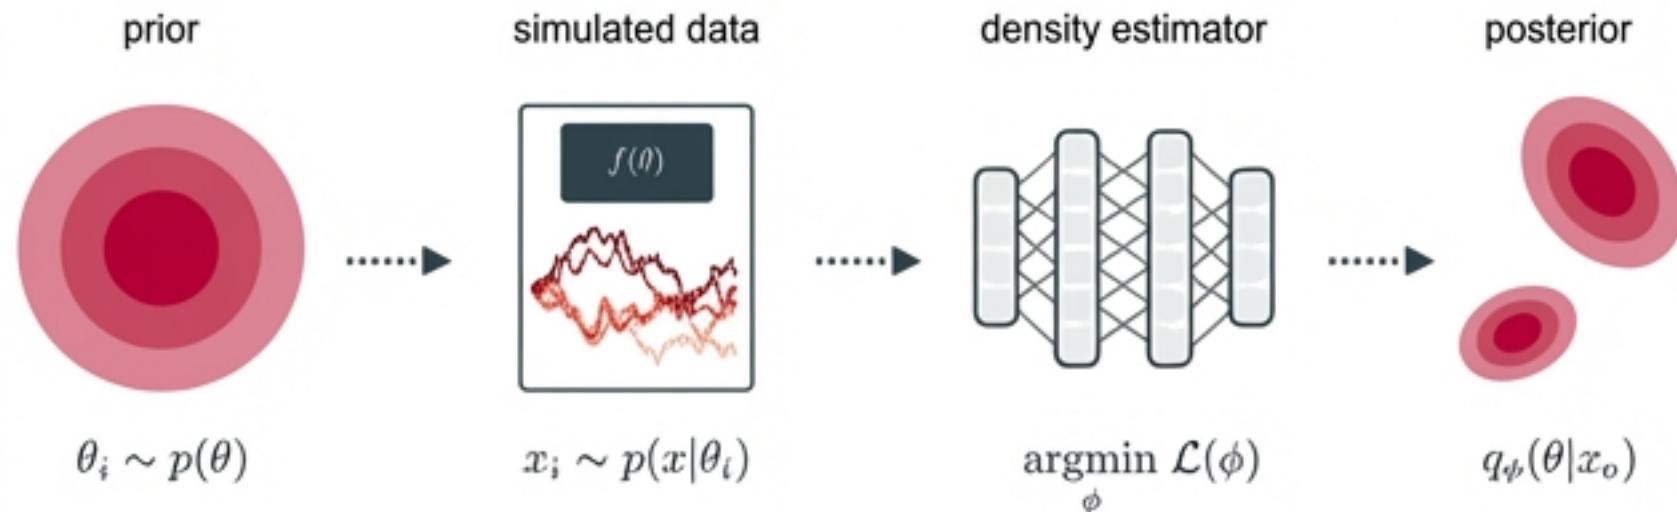
**FOCUS**

Architecture & Roadmap

> FROM RESEARCH CODE TO USER-ORIENTED PACKAGE

# High-Level Organization: Mapping Science to Source

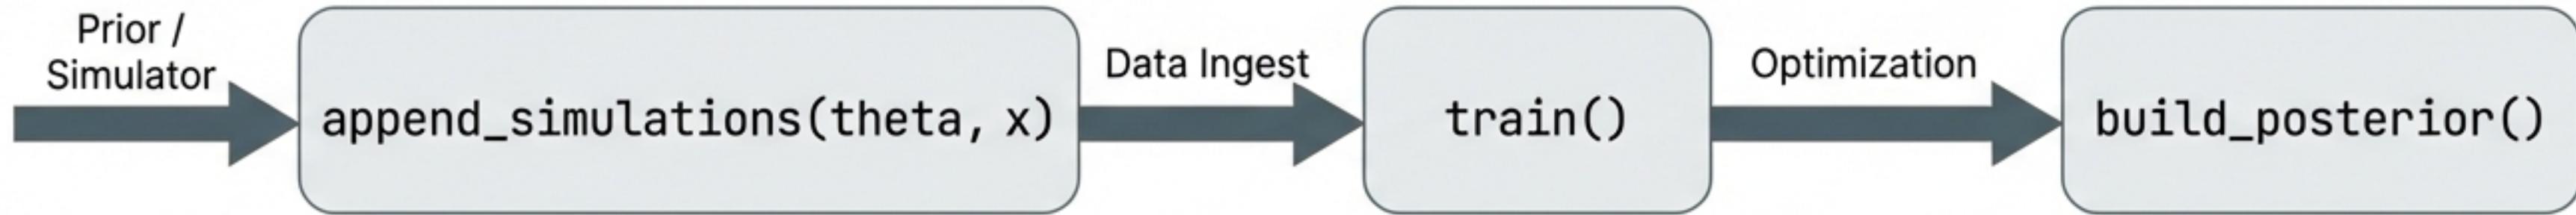
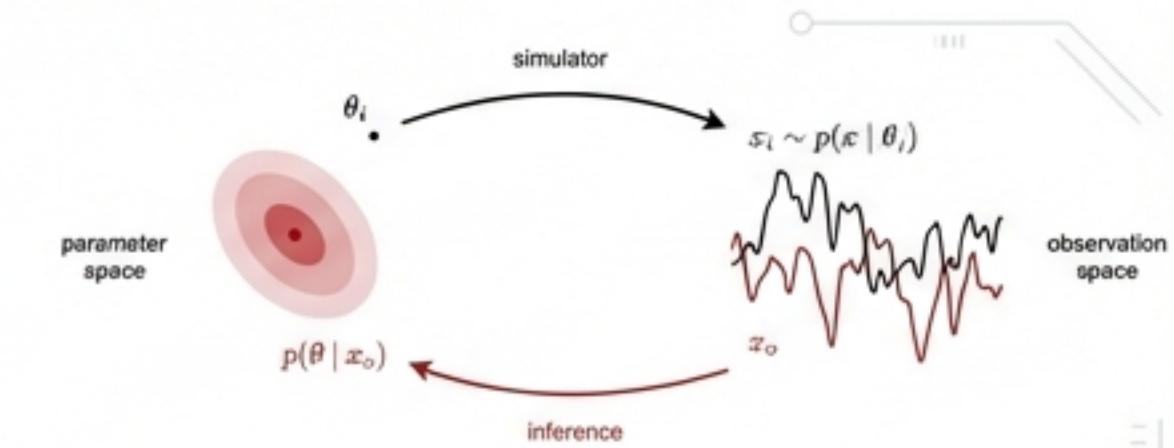
## THE MENTAL MODEL (Science)



## THE REPOSITORY (Code)

```
sbi/
  ├── inference/   <-- The "Brains" (Trainers, Algorithms)
  │   ├── trainers/
  │   └── posteriors/
  ├── neural_nets/ <-- The "Engine" (Density Estimators)
  │   ├── flow/
  │   └── embedding_nets/
  ├── samplers/    <-- Support (MCMC wrappers)
  ├── utils/       <-- Support (Metrics, user diagnostics)
  └── diagnostics/
```

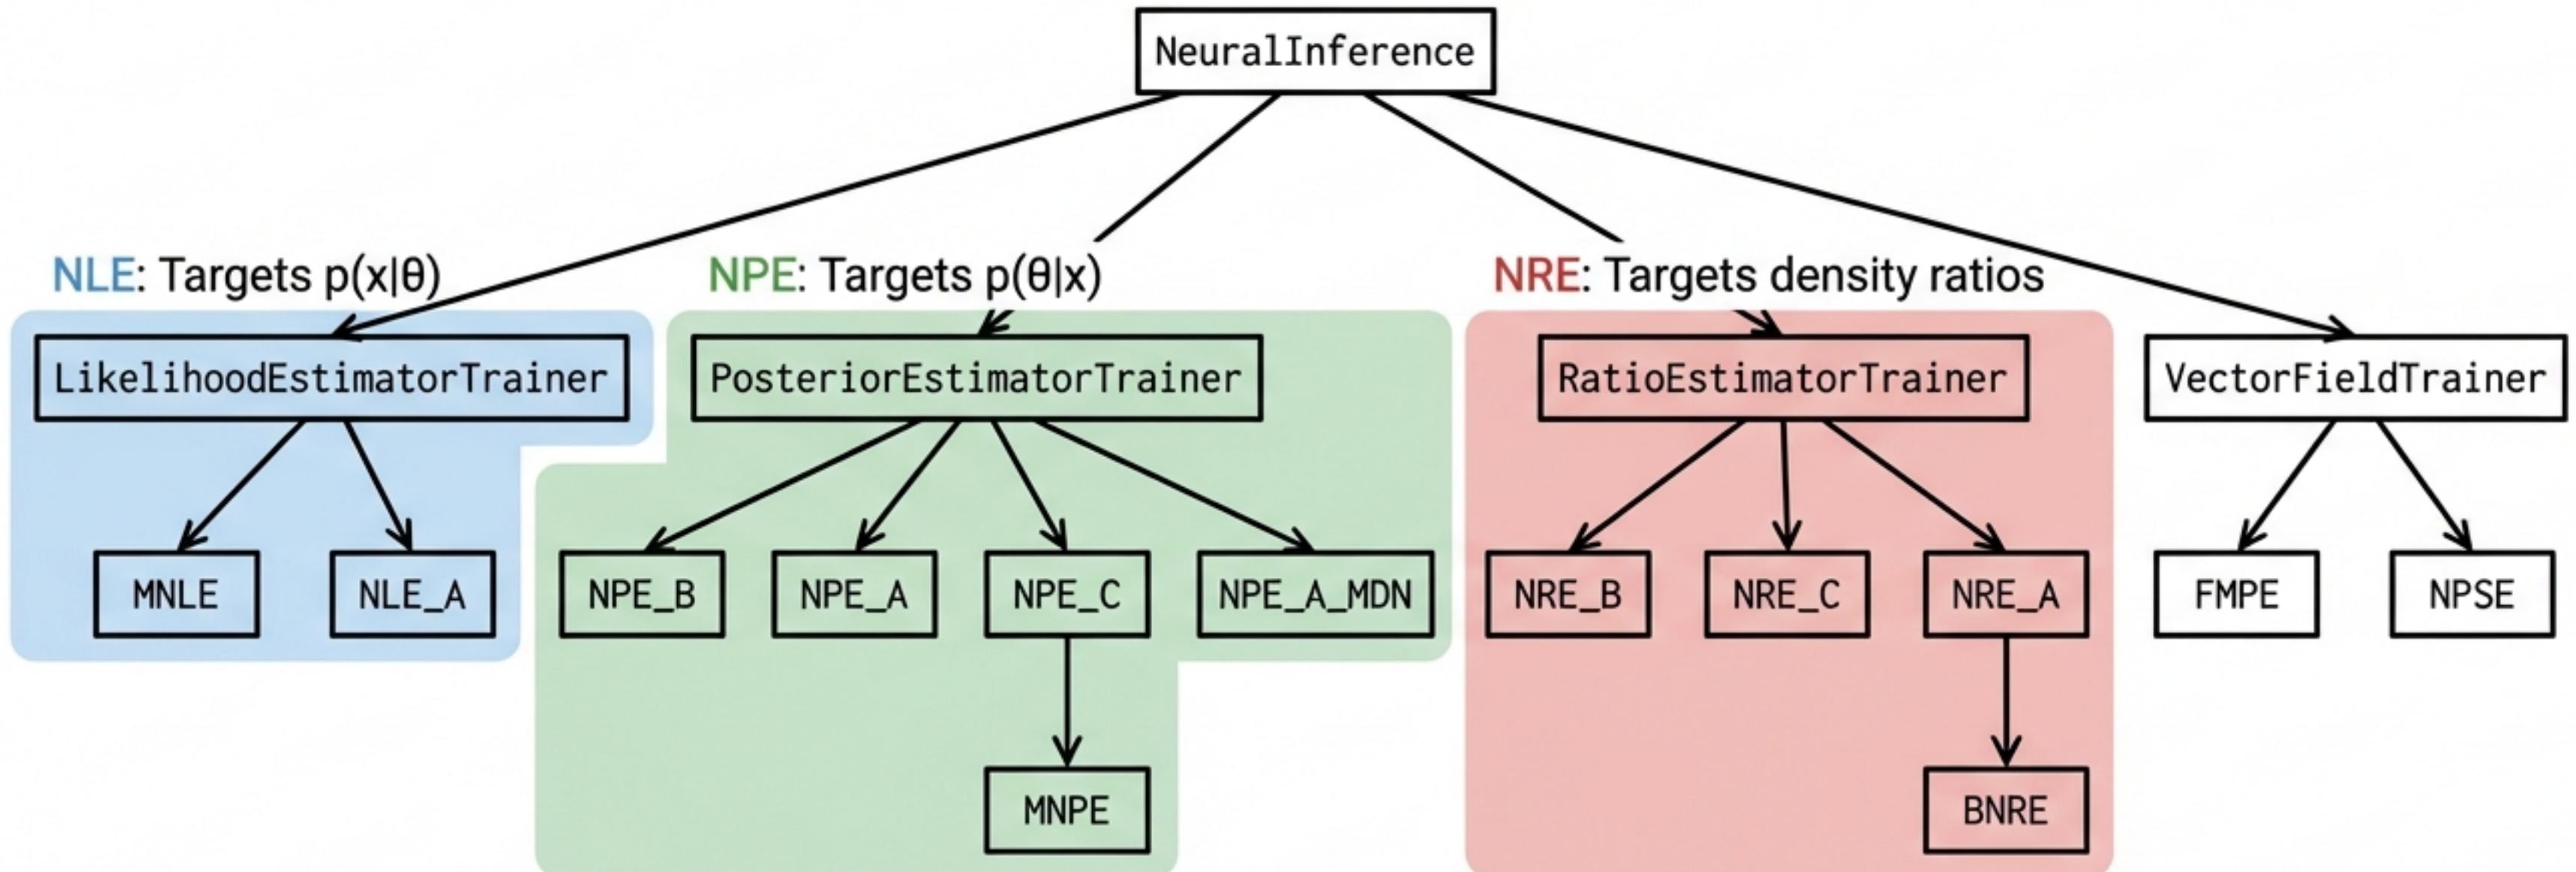
# The Life of an Inference Task



## Design Pattern: Trainer-Posterior Separation

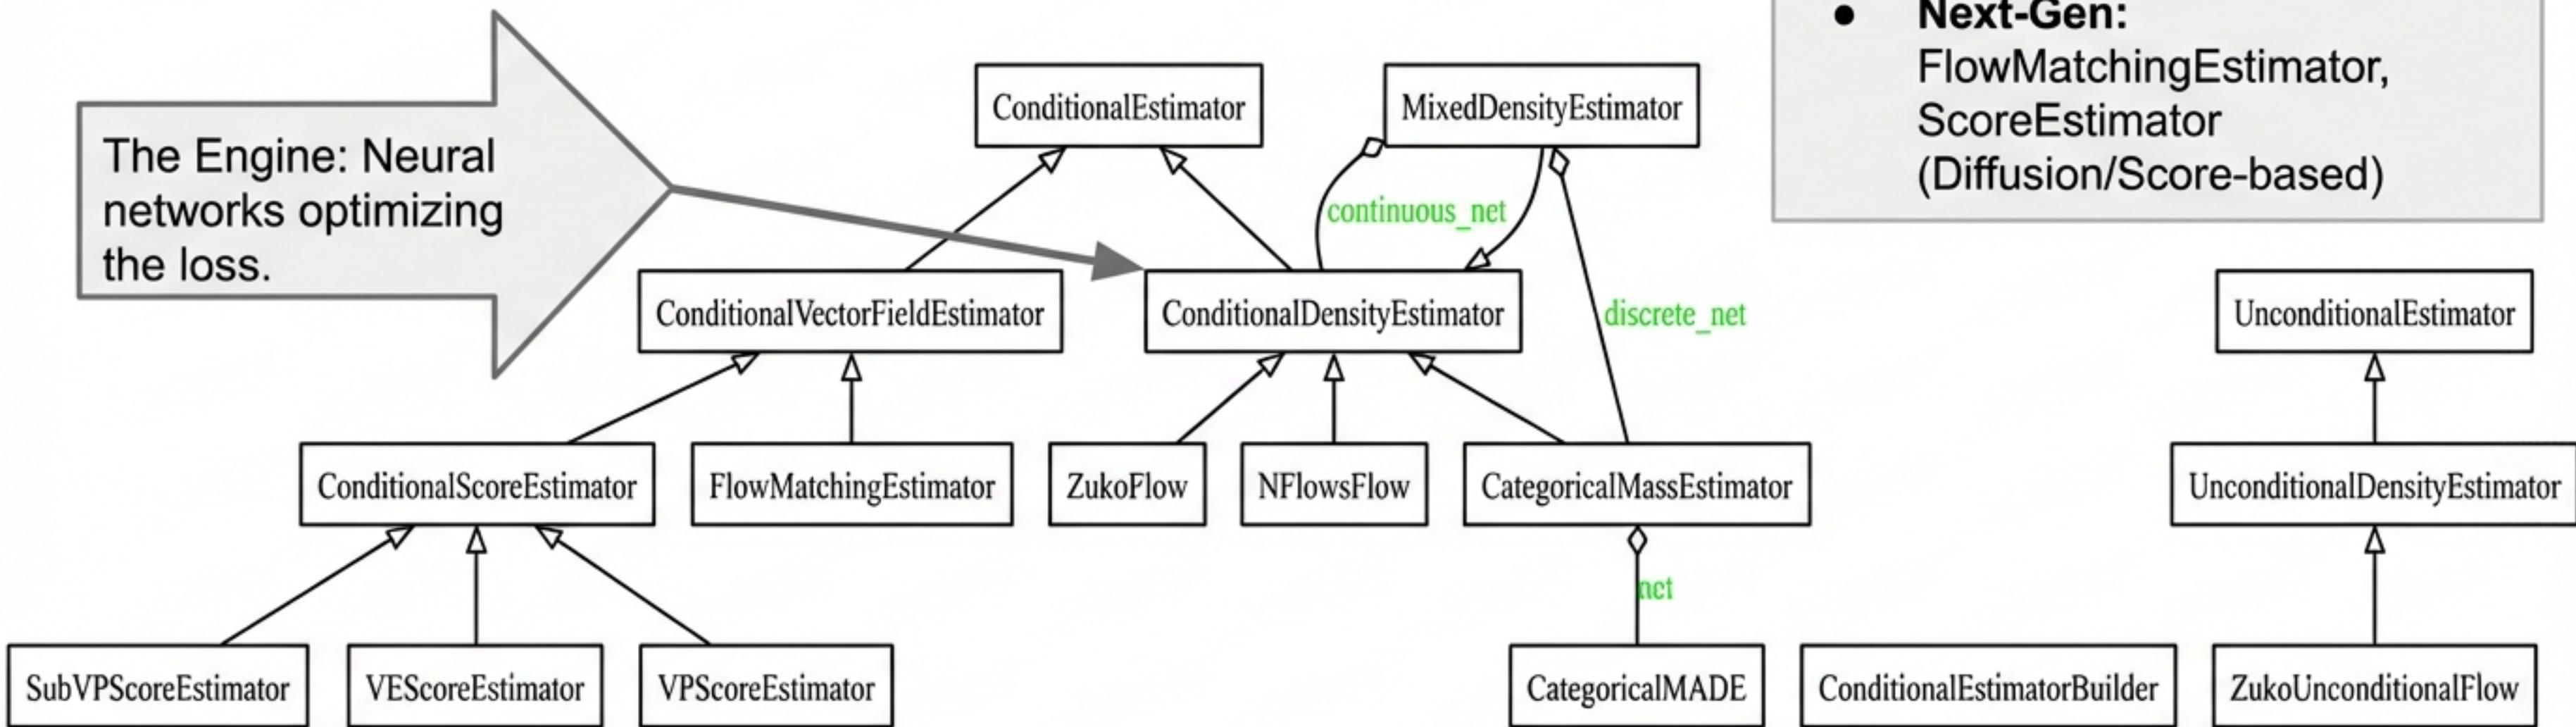
- **Trainers:** Handle optimization loop and simulation management.
- **Posteriors:** Handle sampling, log\_prob evaluation, and device management.
- **Why?** Allows swapping sampling backends (e.g., Rejection vs. MCMC) after training without retraining the network.

# Orchestrating Inference: The Trainer Hierarchy



Root: **NeuralInference** handles the common training loop. Leaf nodes (e.g., **NPE\_C**) represent specific algorithm variants.

# Powering Inference: Density Estimators



Evolution Note: Moving away from factory functions (`posterior_nn`) toward direct class instantiation (See [Issue #16501](#))

# The User Interface: Posterior Classes



## Direct / Amortized

- **Class:** `DirectPosterior`
- **Use Case:** NPE methods
- **Feature:** Instant sampling via the normalizing flow's inverse pass.

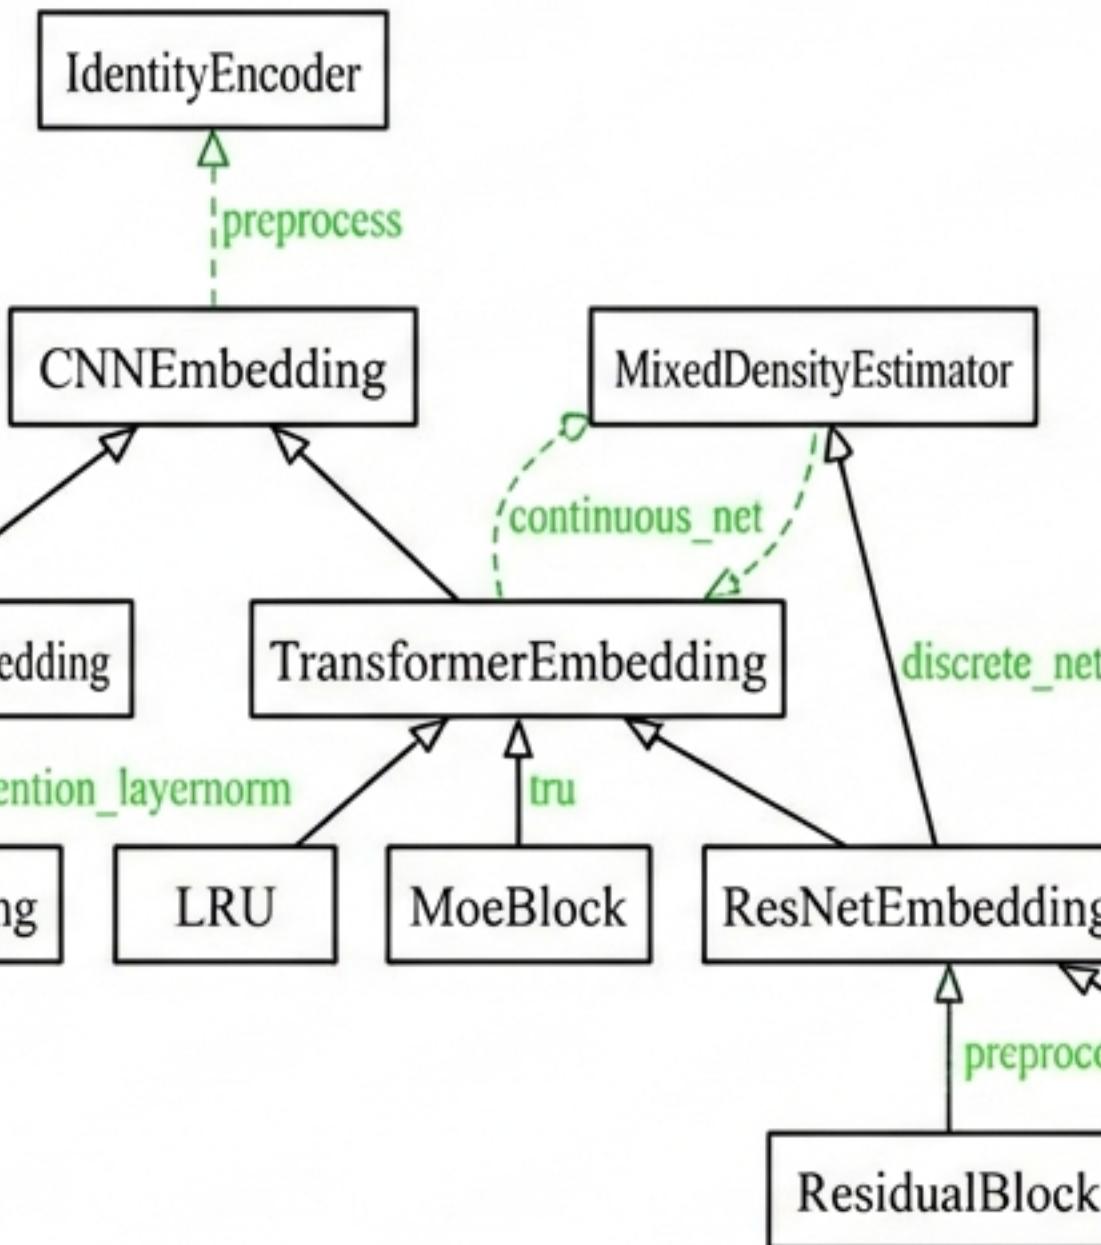
## Sampling-Based (MCMC/VI)

- **Classes:** `MCMCPosterior`, `VTPosterior`, `RejectionPosterior`
- **Use Case:** NLE / NRE methods
- **Mechanism:** Potential Functions  
 $\text{Potential} = \log_{\text{likelihood}} + \log_{\text{prior}}$   
(Allows flexible sampler choices via `sbi.samplers`)

# Handling Complex Data: Embedding Networks

## Role:

Pre-process high-dimensional data (images, time-series) into feature vectors before they hit the density estimator.



## Standard Types

- CNNEmbedding (Images)
- ResNetEmbedding (Deep features)
- TransformerEmbedding (Sequential data)
- PermutationInvariantEmbedding (Sets/IID data)

## Implementation Note

Must inherit from `nn.Module` and is trained end-to-end with the inference network.

# The API Contract: Essential Methods

Trainer Contract (sbi.inference)

```
def append_simulations(theta, x): # Add training data  
def train(): # Execute training loop  
def build_posterior(): # Return posterior object
```

Posterior Contract (sbi.inference.posteriors)

```
def sample(shape, x): # Generate parameters  
def log_prob(theta, x): # Evaluate density  
def map(x): # Find Maximum A Posteriori
```

Density Estimator Contract (sbi.neural\_nets)

```
def loss(inputs, condition): # Calculate training objective
```



sbi-dev / sbi

Type / to search

&lt; Code

# Development Setup: Getting Ready to Code

## Part 2



sbi

Public

Edit Pins

Unwatch

25



```
$ git clone https://github.com/username/sbi.git
$ cd sbi

$ # CRITICAL: Install in editable mode with dev dependencies
$ uv pip install -e ".[dev]"

# Install hooks to catch style issues early
$ pre-commit install
```

docs: prepare release (#1447)

1,304 Commits

6 months ago

4 days ago

2 hours ago

39 minutes ago

39 minutes ago

4 days ago

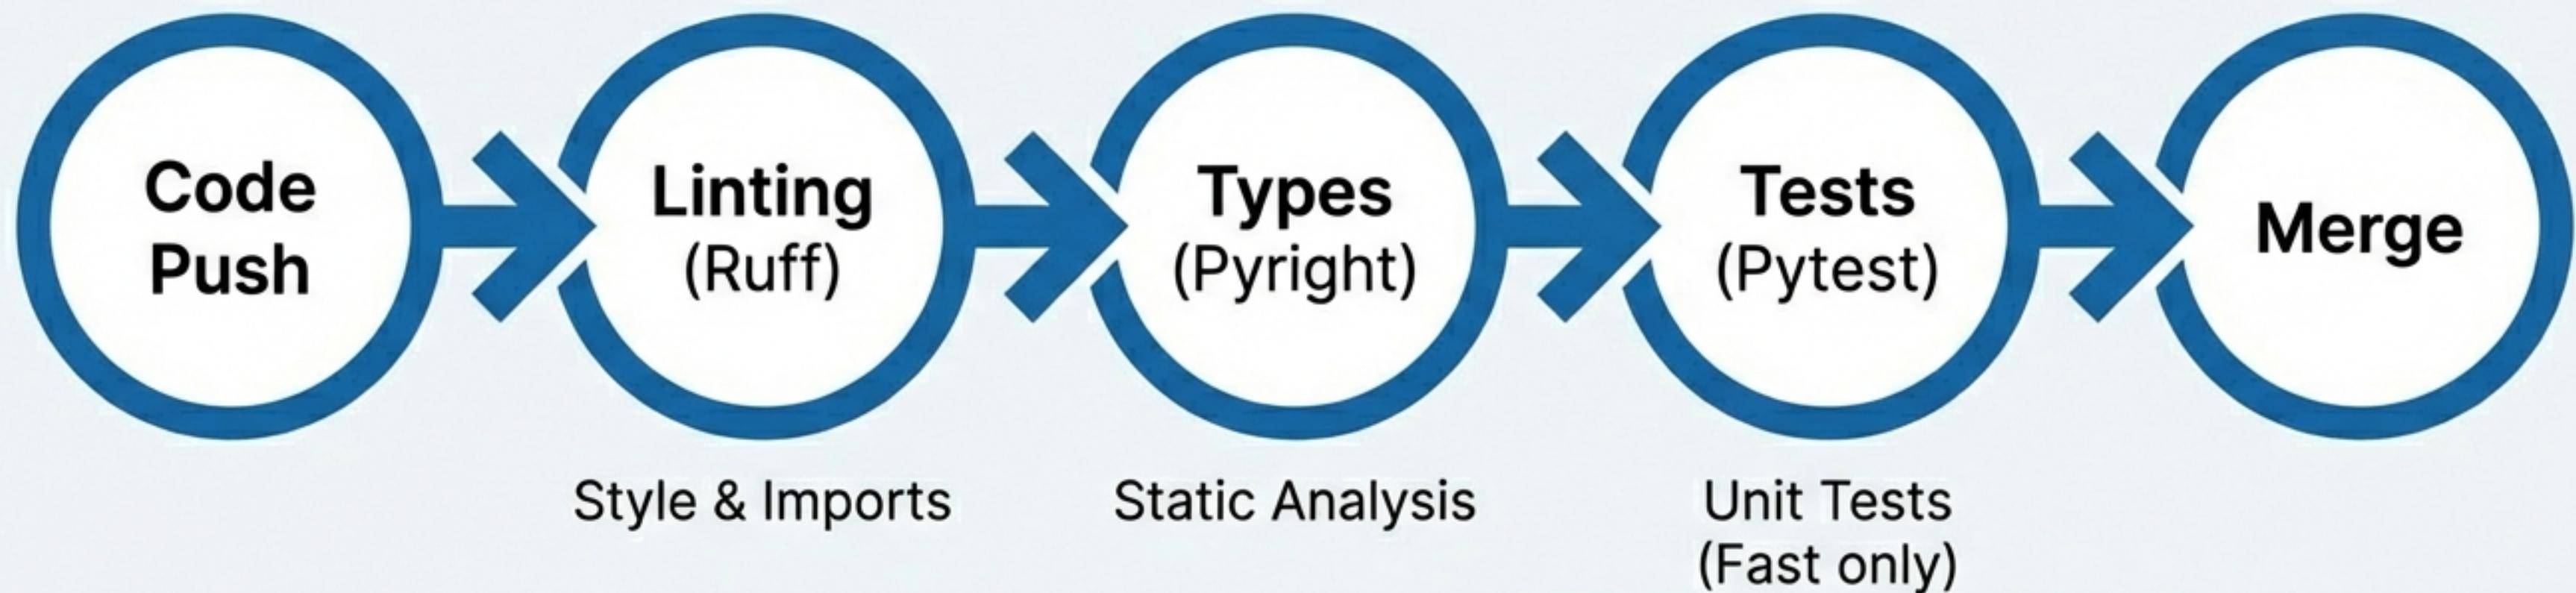
**Pro Tip:** We use 'uv' for modern, fast Python package management.



build: drop py3.9, fix deps, fix types, lc2st tests (#1412)

NotebookLM

# The CI/CD Pipeline: What Happens When You Push?



**Note:** CI runs fast tests only.  
GPU and slow tests are  
separate workflows.



sbi-dev / sbi

Type / to search

Code

Pull requests



# Testing Best Practices



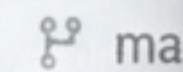
sbi

Public

Edit Pins

Unwatch

25



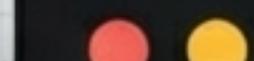
main



3



.dev



pytest -m "not slow and not gpu"

file

Code

6 ago

1,304 Commits

6 months ago

## Structure

Tests live in tests/. Use fixtures for common setup to reduce runtime.

## Markers

Decorate heavy tests with:  
`@pytest.mark.slow`  
`@pytest.mark.gpu`

## Performance

Check Mini-SBIBM for benchmarking new methods against standard baselines.

# PR Etiquette: Getting Merged

## Title

feat: Add flow matching estimator

## Description

1. WHAT does it do?
2. WHY is it needed?
3. HOW does it work?
4. Fixes Issue #...

## Pre-Flight Checklist

- Passes ruff & pyright
- Tests pass locally
- New functionality has new tests
- Documentation updated

# Pick Your Challenge: Hackathon Projects

## Good First Issues

- Docstring cleanup (#1534)

- Expose FMPE/NPSE hyperparameters (#1284)

- GPU-compatible diagnostics (#1160)

## Documentation & Guides

- Hyperparameter tuning guide (#1346)

- Saving/Loading estimators guide (#1525)

- Pyro integration guide (#1517)

## Advanced / Architecture

- Better flow matching training (#1445)

- Refactor NPE\_A\_MDN (#1408)

- Remove posterior\_nn factory functions (#1659)

# Ready? Let's Build.



## Next Steps

1. 60 min interactive discussion
2. Form working groups
3. Assign mentors